

Chapter 6:

Physical Database Design and Performance

Objectives

- Definition of terms
- Describe the physical database design process
- Choose storage formats for attributes
- Select appropriate file organizations
- Describe three types of file organization
- Describe indexes and their appropriate use
- Translate a database model into efficient structures
- Know when and how to use denormalization

Physical Database Design

- Purpose—translate the logical description of data into the *technical specifications* for storing and retrieving data
- Goal—create a design for storing data that will provide *adequate performance* and insure *database integrity, security, and recoverability*

Physical Design Process

Inputs

- Normalized relations
- Volume estimates
- Attribute definitions
- Response time expectations
- Data security needs
- Backup/recovery needs
- Integrity expectations
- DBMS technology used



Leads to

Decisions

- Attribute data types
- Physical record descriptions
(doesn't always match logical design)
- File organizations
- Indexes and database architectures
- Query optimization

Figure 6-1 Composite usage map

(Pine Valley Furniture Company)

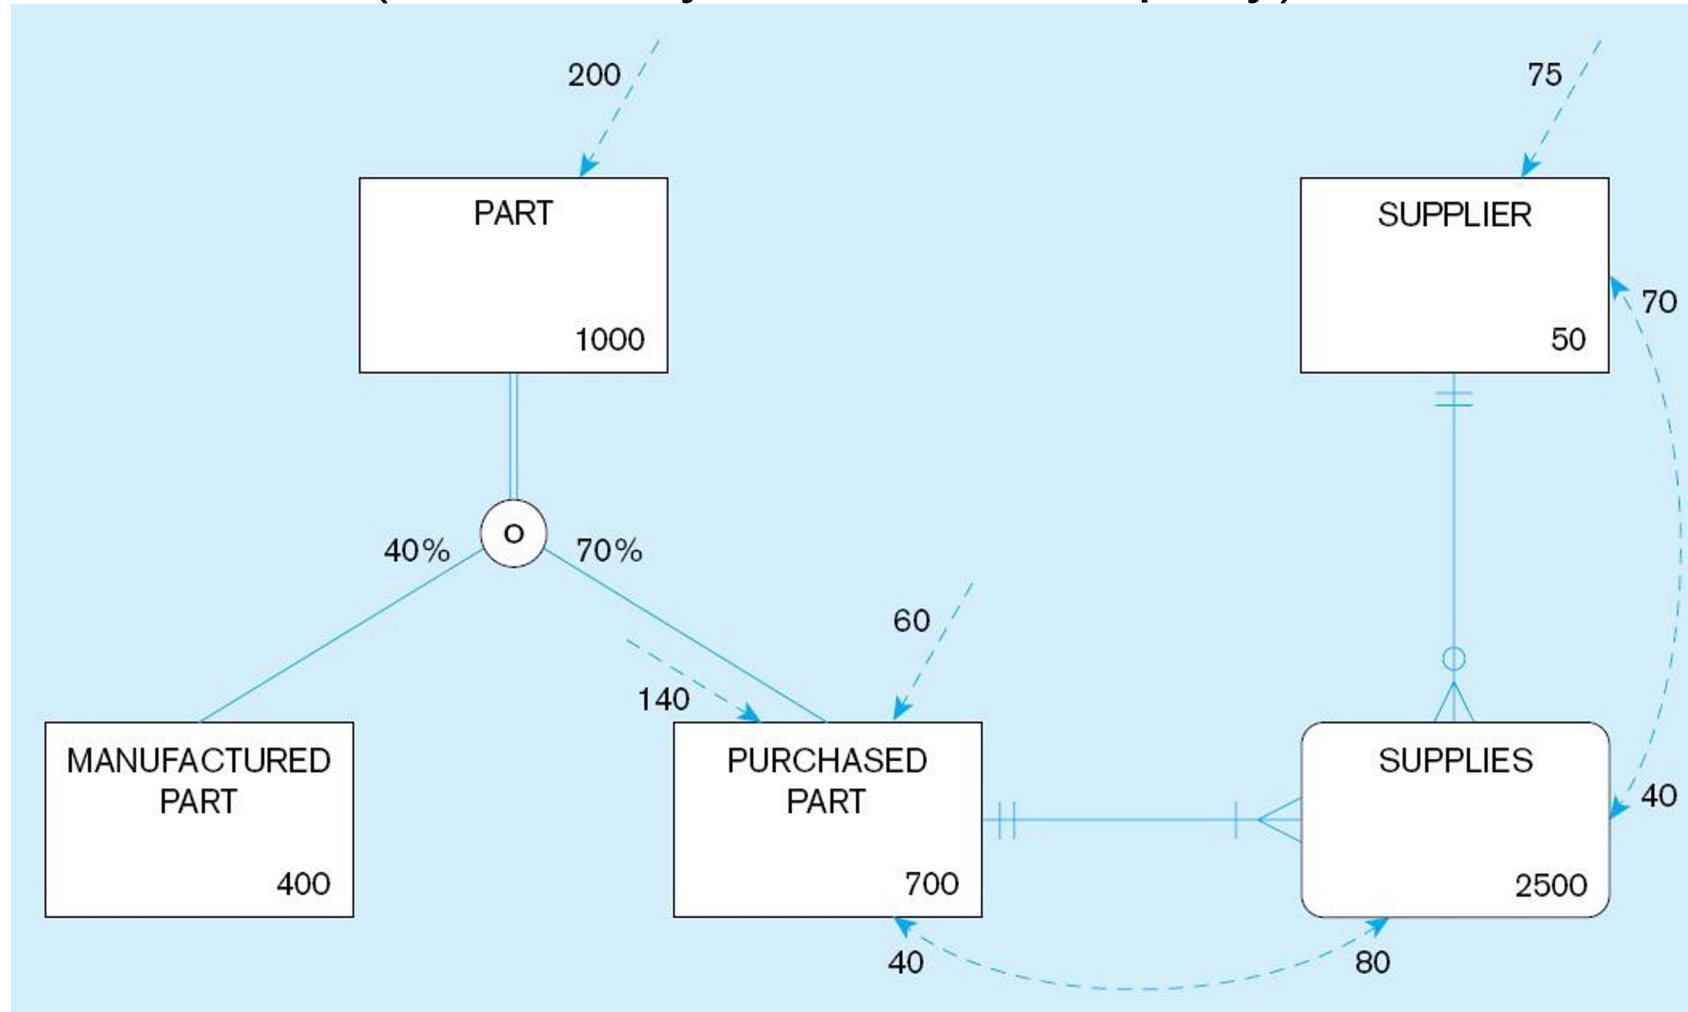


Figure 6-1 Composite usage map

(Pine Valley Furniture Company) (cont.)

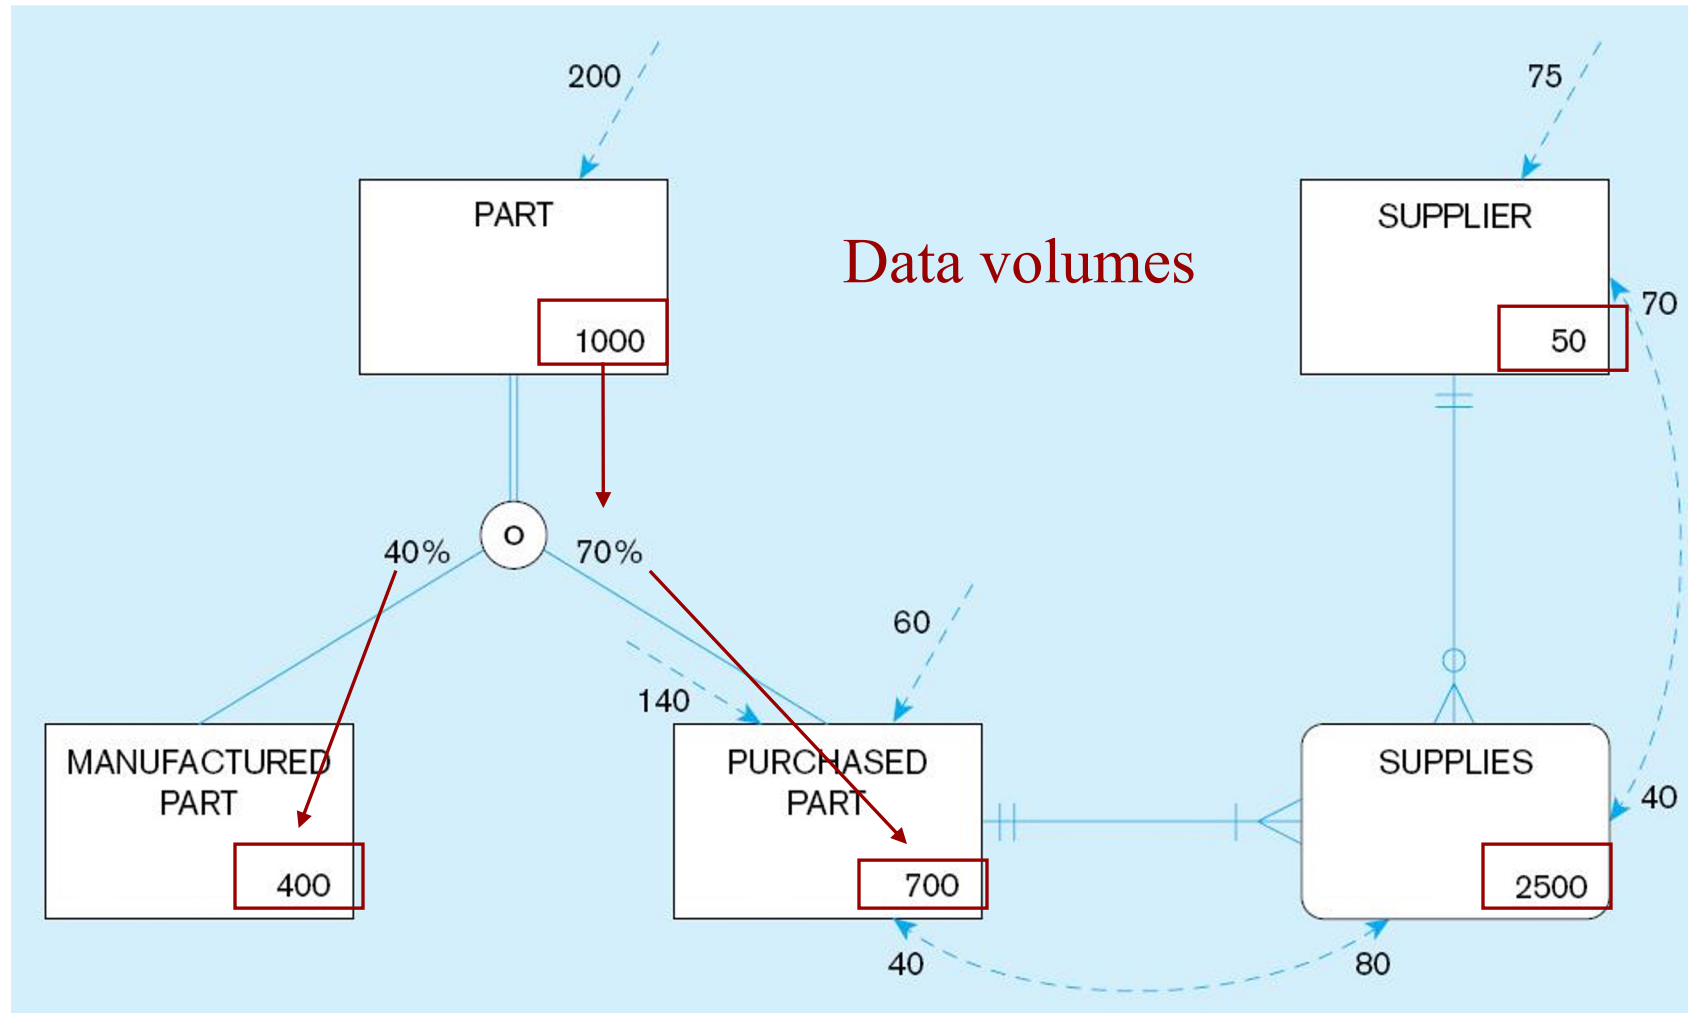


Figure 6-1 Composite usage map

(Pine Valley Furniture Company) (cont.)

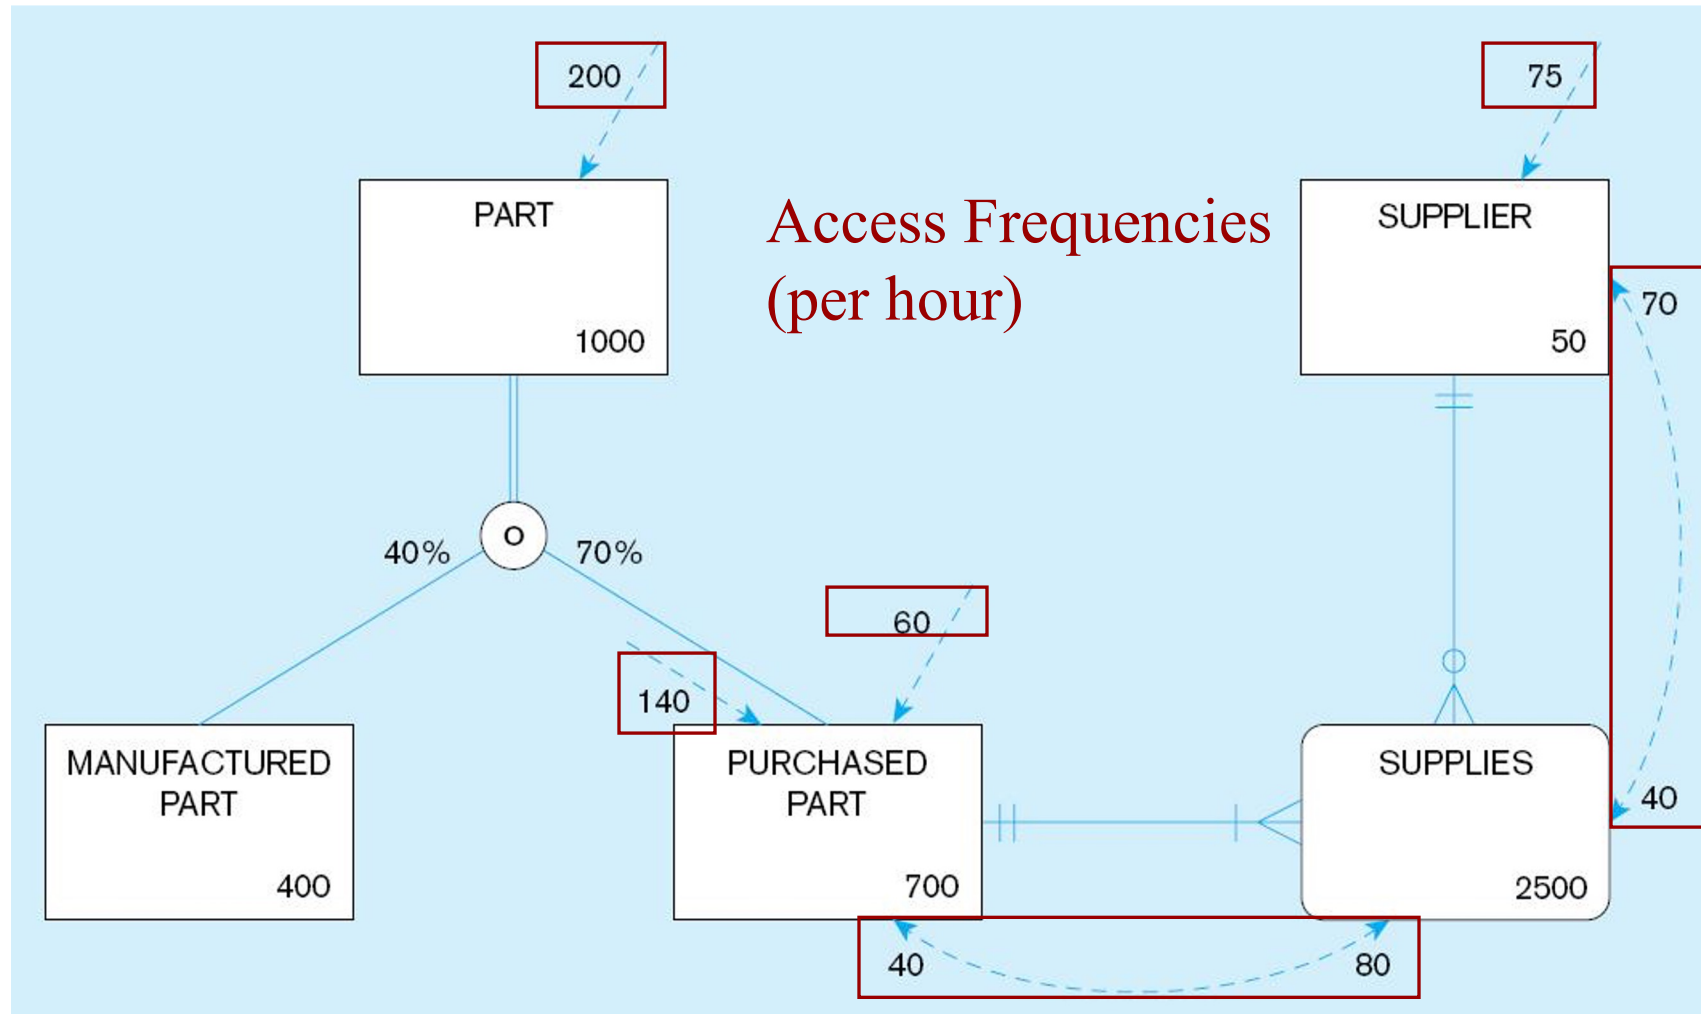


Figure 6-1 Composite usage map

(Pine Valley Furniture Company) (cont.)

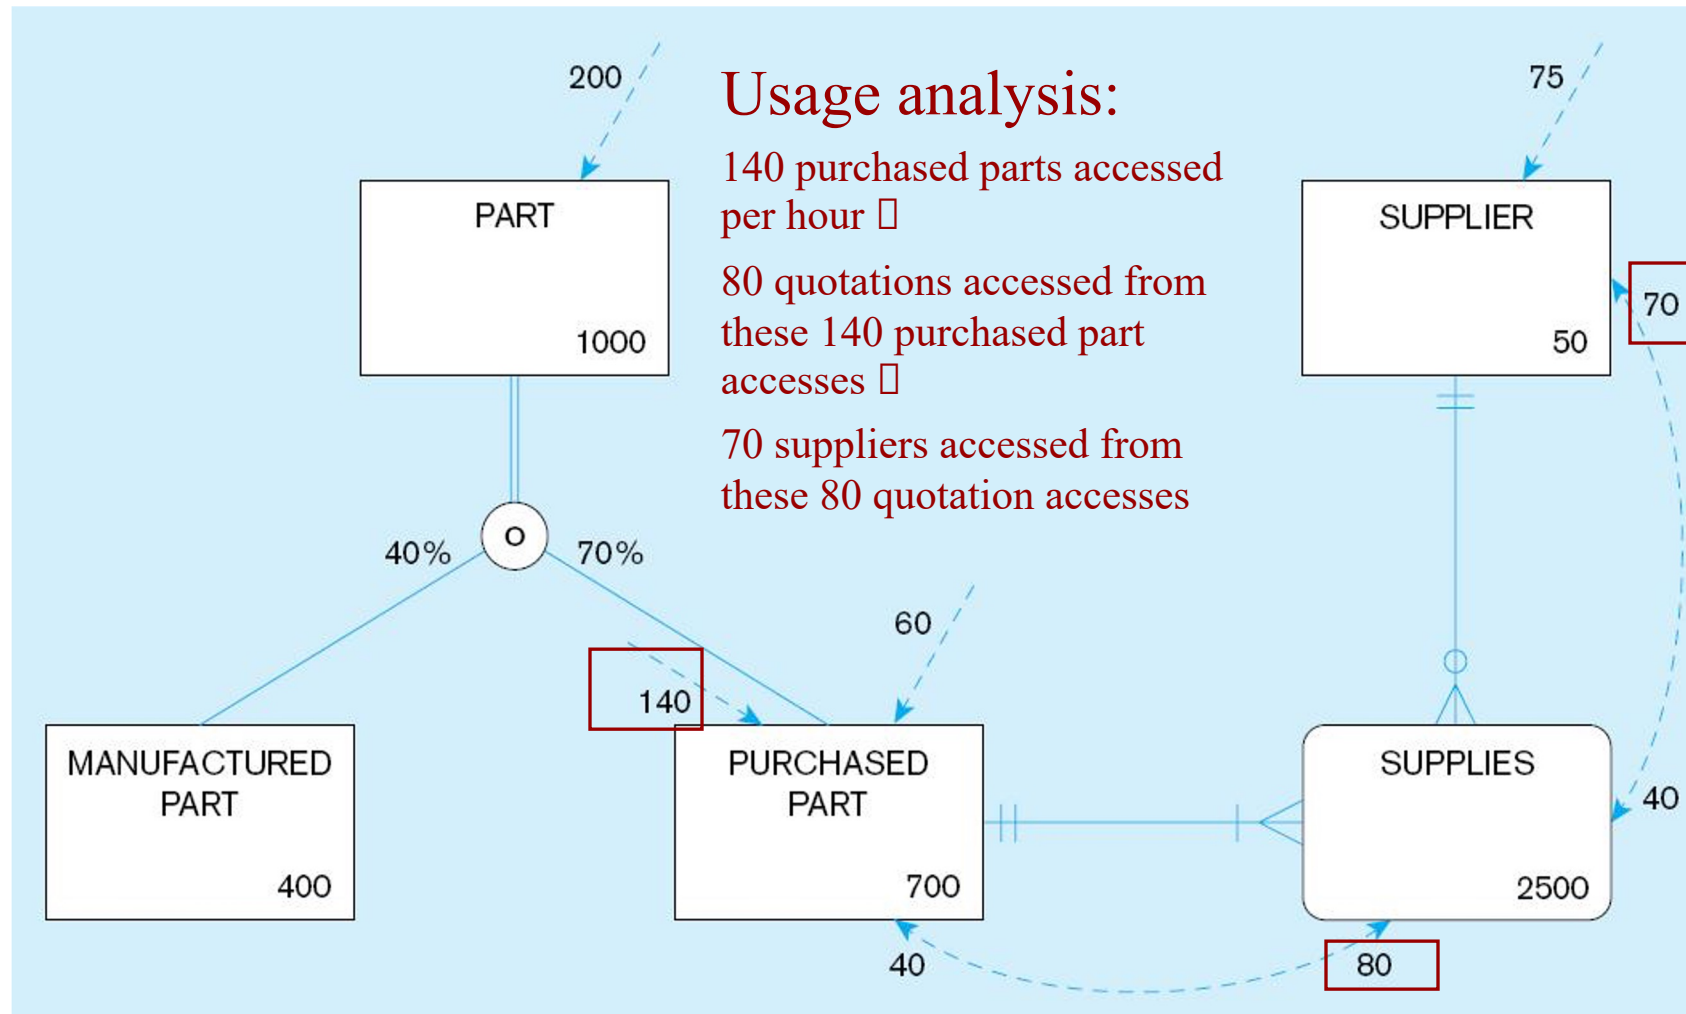
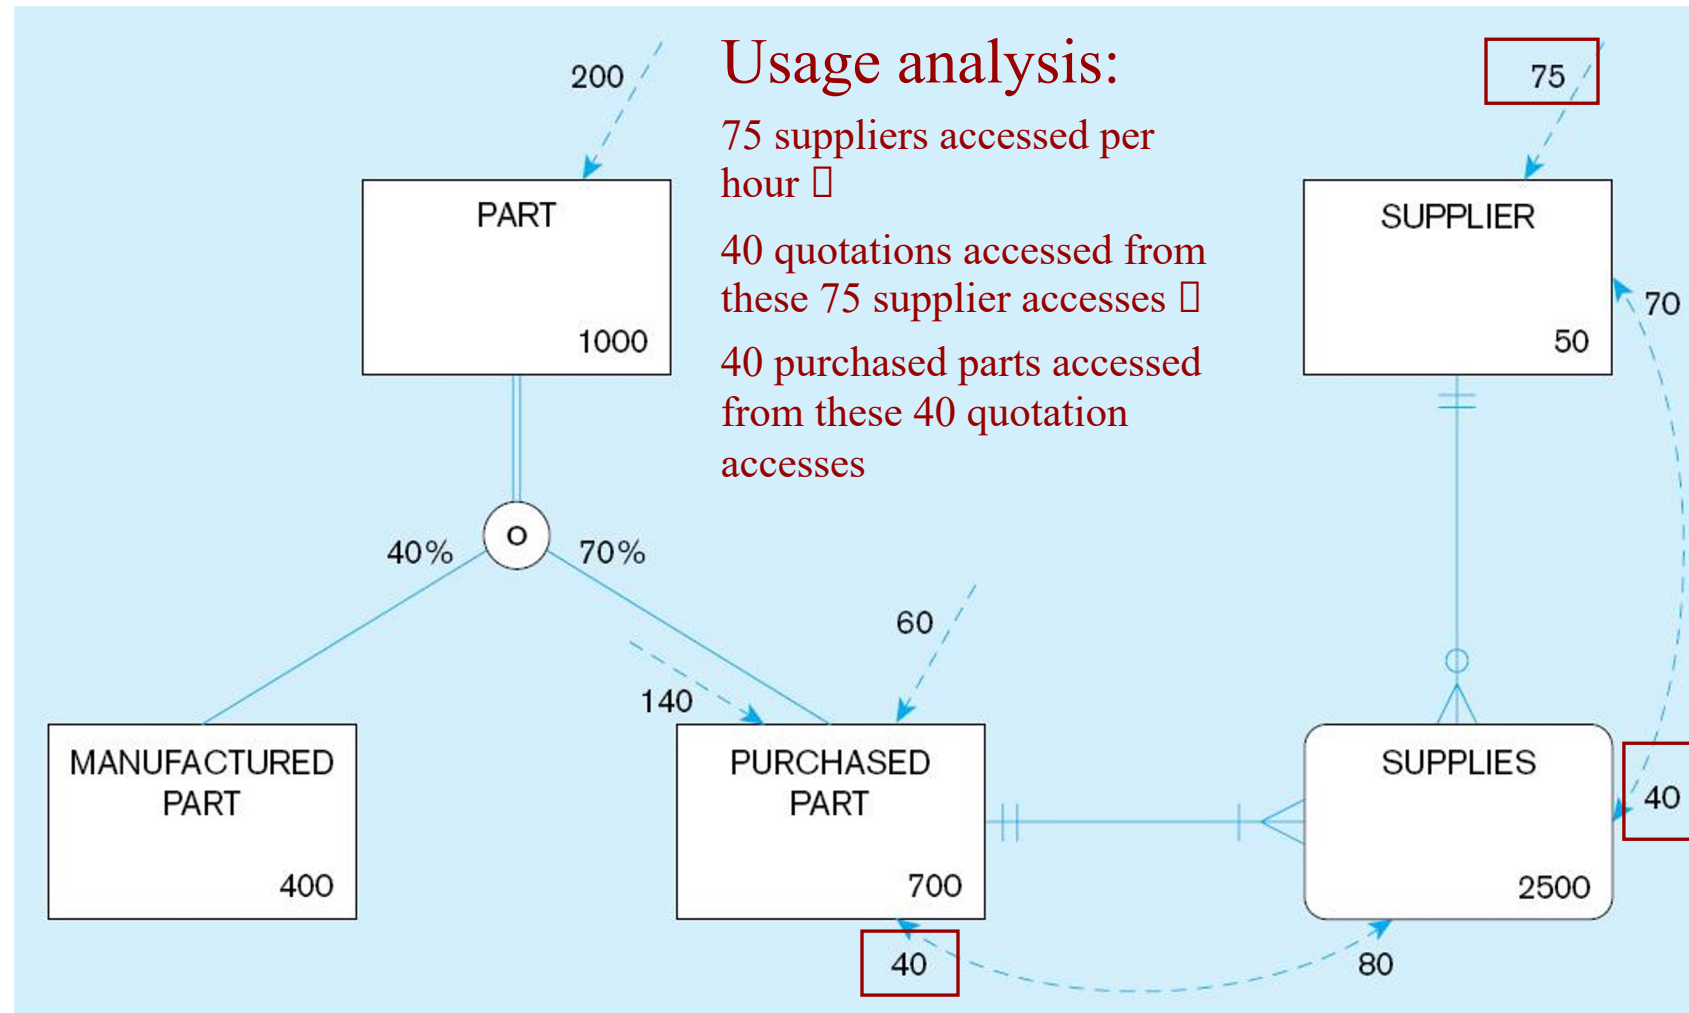


Figure 6-1 Composite usage map

(Pine Valley Furniture Company) (cont.)



Designing Fields

- Field: smallest unit of data in database
- Field design
 - Choosing data type
 - Coding, compression, encryption
 - Controlling data integrity

Choosing Data Types

- CHAR—fixed-length character
- VARCHAR2—variable-length character (memo)
- LONG—large number
- NUMBER—positive/negative number
- INEGER—positive/negative whole number
- DATE—actual date
- BLOB—binary large object (good for graphics, sound clips, etc.)

Figure 6-2 Example code look-up table

(Pine Valley Furniture Company)

PRODUCT File

Product_No	Description	Finish	...
B100	Chair	C	
B120	Desk	A	
M128	Table	C	
T100	Bookcase	B	
...	

FINISH Look-up Table

Code	Value
A	Birch
B	Maple
C	Oak

Code saves space, but costs an additional lookup to obtain actual value

Field Data Integrity

- Default value—assumed value if no explicit value
- Range control—allowable value limitations (constraints or validation rules)
- Null value control—allowing or prohibiting empty fields
- Referential integrity—range control (and null value allowances) for foreign-key to primary-key match-ups

Sarbanes-Oxley Act (SOX) legislates importance of financial data integrity

Handling Missing Data

- Substitute an estimate of the missing value (e.g., using a formula)
- Construct a report listing missing values
- In programs, ignore missing data unless the value is significant (sensitivity testing)

Physical Records

- Physical Record: A group of fields stored in adjacent memory locations and retrieved together as a unit
- Page: The amount of data read or written in one I/O operation

Denormalization

- Transforming ***normalized*** relations into ***unnormalized*** physical record specifications
- Benefits:
 - Can improve performance (speed) by reducing number of table lookups (i.e. *reduce number of necessary join queries*)
- Costs (due to data duplication)
 - Wasted storage space
 - Data integrity/consistency threats
- Common denormalization opportunities
 - One-to-one relationship (Fig. 6-3)
 - Many-to-many relationship with attributes (Fig. 6-4)
 - Reference data (1:N relationship where 1-side has data not used in any other relationship) (Fig. 6-5)

Figure 6-3 A possible denormalization situation: two entities with one-to-one relationship

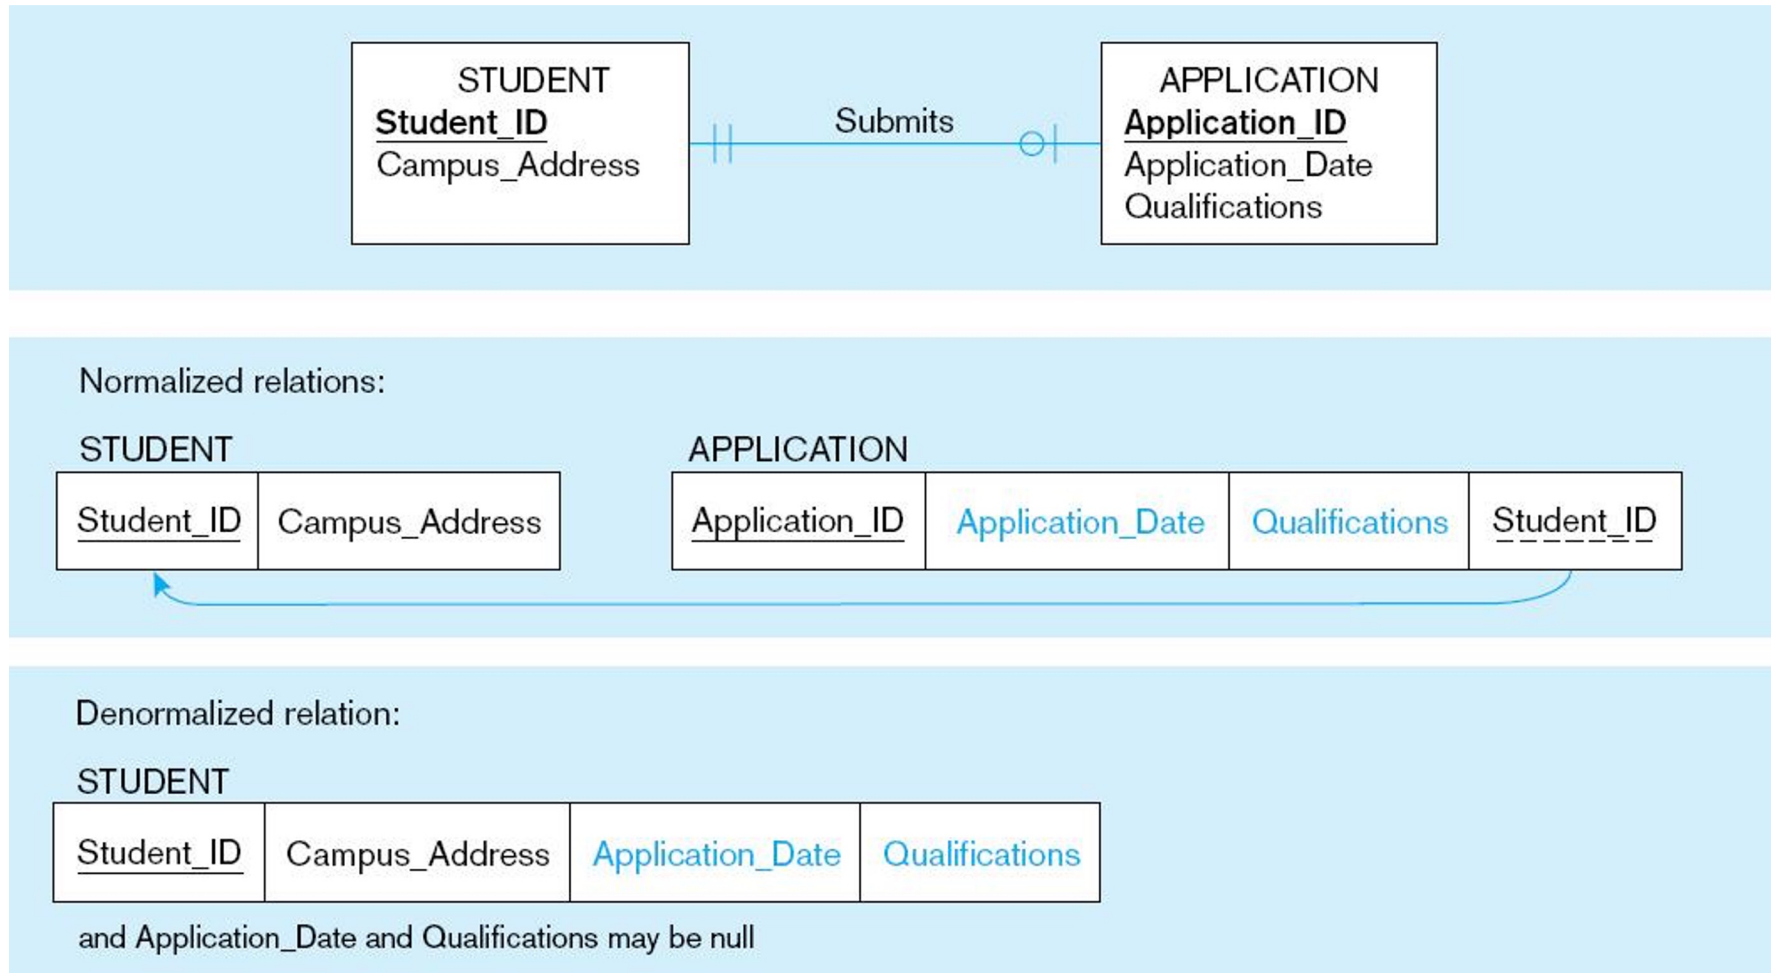


Figure 6-4 A possible denormalization situation: a many-to-many relationship with nonkey attributes

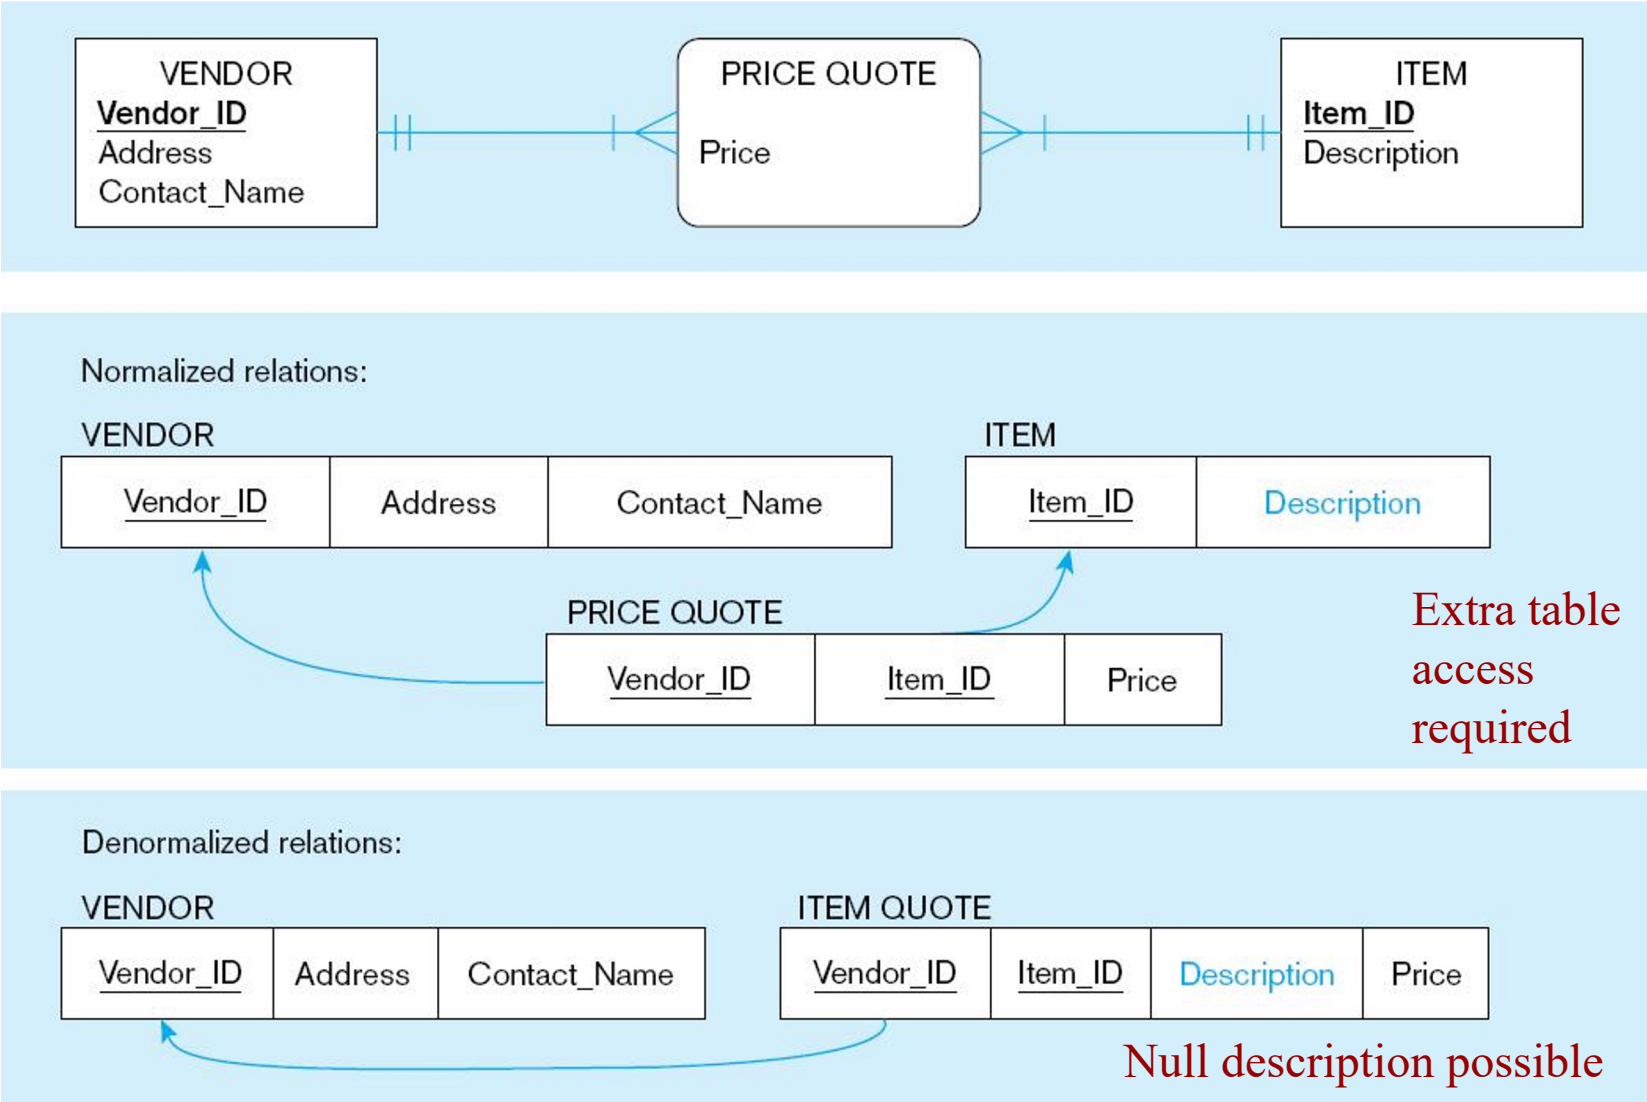


Figure 6-5
A possible
denormalization
situation:
reference data



Normalized relations:

STORAGE

<u>Instr_ID</u>	Where_Store	Container_Type
-----------------	-------------	----------------

ITEM

<u>Item_ID</u>	Description	Instr_ID
----------------	-------------	----------

Extra table
access
required

Denormalized relation:

ITEM

<u>Item_ID</u>	Description	Where_Store	Container_Type
----------------	-------------	-------------	----------------

Data duplication

Partitioning

- Horizontal Partitioning: Distributing the rows of a table into several separate files
 - Useful for situations where different users need access to different rows
- Vertical Partitioning: Distributing the columns of a table into several separate relations
 - Useful for situations where different users need access to different columns
 - The primary key must be repeated in each file
- Combinations of Horizontal and Vertical

Partitioning (cont.)

- **Advantages of Partitioning:**
 - Efficiency: Records used together are grouped together
 - Local optimization: Each partition can be optimized for performance
 - Security, recovery
 - Load balancing: Partitions stored on different disks, reduces contention
 - Take advantage of parallel processing capability
- **Disadvantages of Partitioning:**
 - Inconsistent access speed: Slow retrievals across partitions
 - Complexity: Non-transparent partitioning
 - Extra space or update time: Duplicate data; access from multiple partitions

Data Replication

- Purposely storing the same data in multiple locations of the database
- Improves performance by allowing multiple users to access the same data at the same time with minimum contention
- Sacrifices data integrity due to data duplication
- Best for data that is not updated often